

# Implementing High Performance Zero-Knowledge ML Provers

Lilia Tang, Alluri Siddhartha Varma, Siheng Pan, Daniel Kang  
UIUC

# Implementing High Performance Zero-Knowledge ML Provers

Lilia Tang, Alluri Siddhartha Varma, Siheng Pan, Daniel Kang  
UIUC

# Increasing calls for ML transparency

## CDRH Issues Guiding Principles for Transparency of Machine Learning-Enabled Medical Devices

### A call for AI data transparency

April 11, 2024

FOR  
June



### It's time to reveal all recommendation algorithms – by law if necessary

By Wa

We should know why we see what we see, not be left in the dark

A necessary step: verify that ML model outputs are honestly computed

# Increasing calls for ML transparency

**CDRH Issues Guiding Principles for  
Transparency of Machine Learning-Enabled  
Medical Devices**

**A call for AI data transparency**

April 11, 2024

FOR  
June



By Wa

**It's time to reveal all recommendation  
algorithms – by law if necessary**

We should know why we see what we see, not be left in the dark

A necessary step: verify that ML model outputs are honestly computed

# Increasing calls for ML transparency

**CDRH Issues Guiding Principles for  
Transparency of Machine Learning-Enabled  
Medical Devices**

**A call for AI data transparency**

April 11, 2024

FOR  
June



**It's time to reveal all recommendation  
algorithms – by law if necessary**

By Wa

We should know why we see what we see, not be left in the dark

A necessary step: verify that ML model outputs are honestly computed

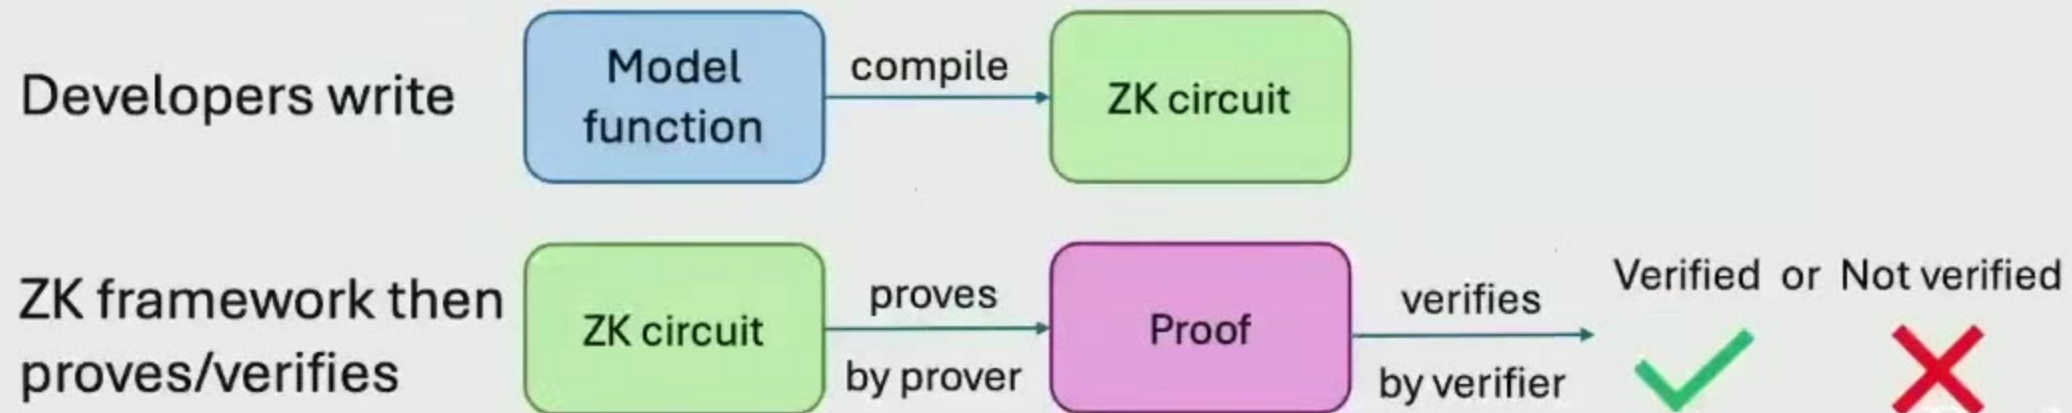




# Zero-Knowledge Succinct Non-interactive Arguments of Knowledge (ZK-SNARK)

A ZK-SNARK is a cryptographic tool enabling a prover to prove a **statement true** to a verifier without disclosing additional information.

E.g., prove that a model is executed correctly without revealing weights and inputs

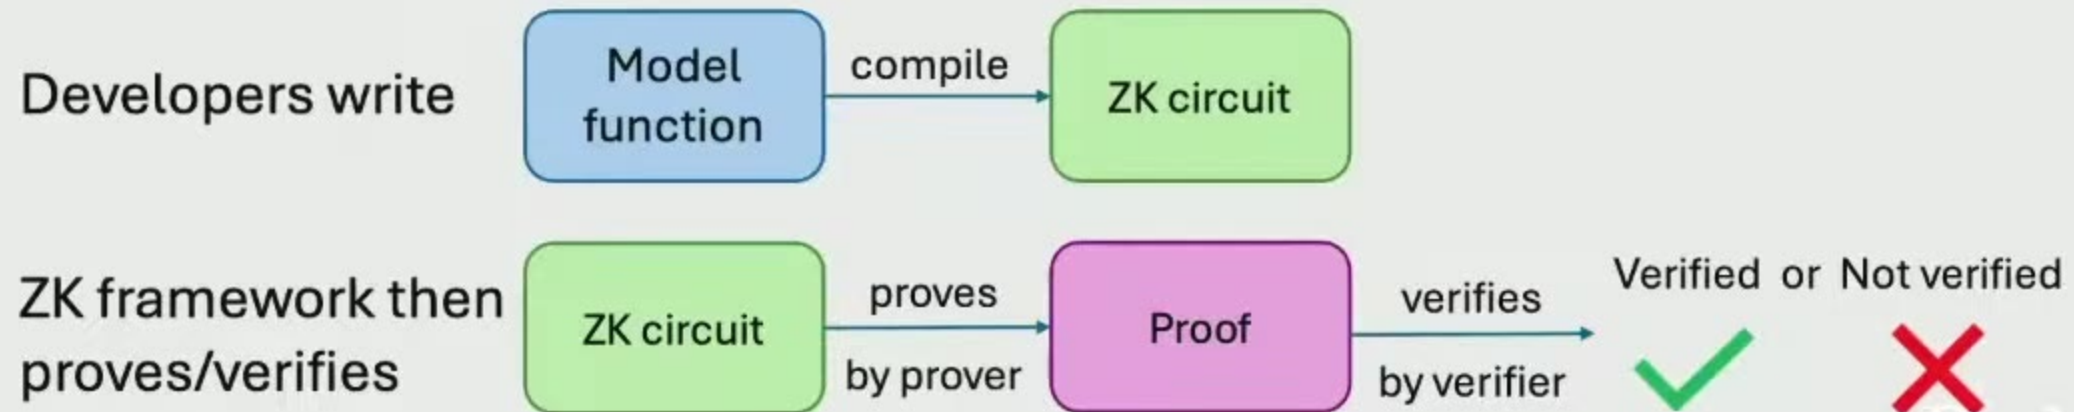


# Zero-Knowledge Succinct Non-interactive Arguments of Knowledge (ZK-SNARK)

Verizon Event ...

A ZK-SNARK is a cryptographic tool enabling a prover to prove a **statement true** to a verifier without disclosing additional information.

E.g., prove that a model is executed correctly without revealing weights and inputs



# Many ZK-SNARK frameworks



**gnark**



**Plonky2**



**==> circom**  
CIRCUIT COMPILER



**halo2**

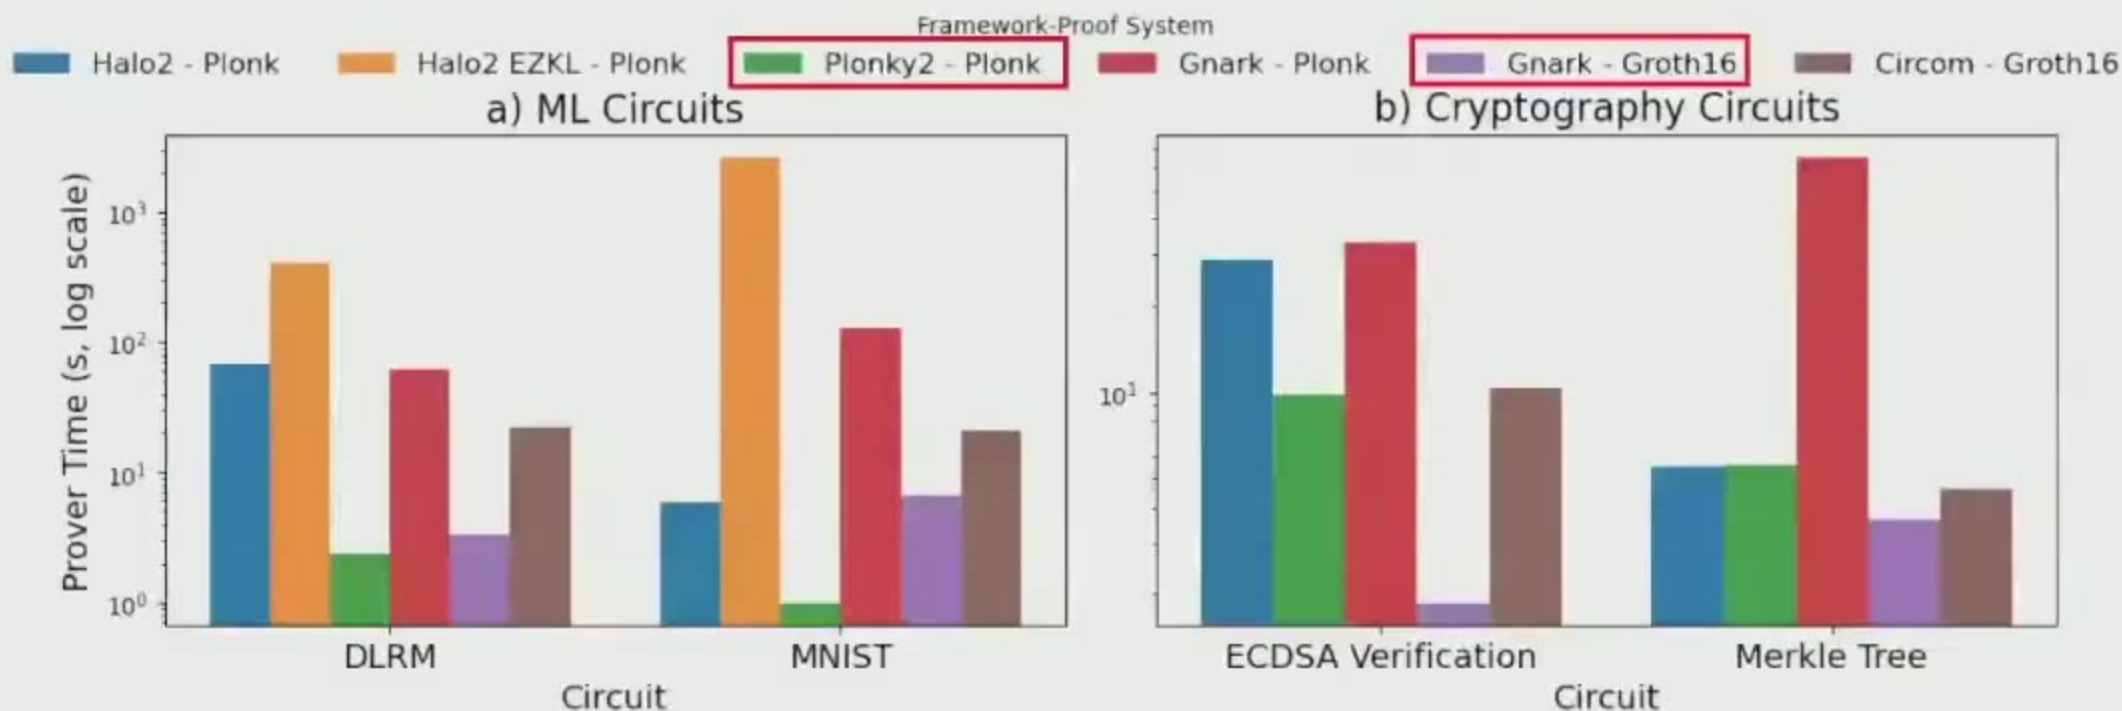
How do we pick which framework to use? And how do we use the framework effectively?



# ZKPerf: a ZK-SNARK proving benchmark

- **ML circuits**
  - **MNIST Convolutional Neural Network:** convolutions
  - **Deep Learning Recommendation Model:** multilayer perceptrons
- **Cryptographic circuits**
  - **Merkle Tree Membership Verification:** hashing
  - **ECDSA:** signatures over elliptic curves
- **Frameworks:** circom (rapidsnark), Halo2, gnark, Plonky2

# Proving time for all tasks and frameworks



- Gnark-Groth16 fastest on crypto circuits but slower on ML
- Plonky2 fastest on ML circuits

# Lookups are key to scalable ML circuits

- Lookups can store input-output mappings for non-linearities
  - Use when expensive to compute with arithmetic constraints
- E.g., ReLU can be proven with bit-decomposition or lookups

$$\text{ReLU}(b) = \max(0, b)$$

ReLU( $b$ ) bit decomposition  
**each operation:**

1. Prove

$$b = \sum_i b_i 2^{l-i} \quad b_i(1 - b_i) = 0$$

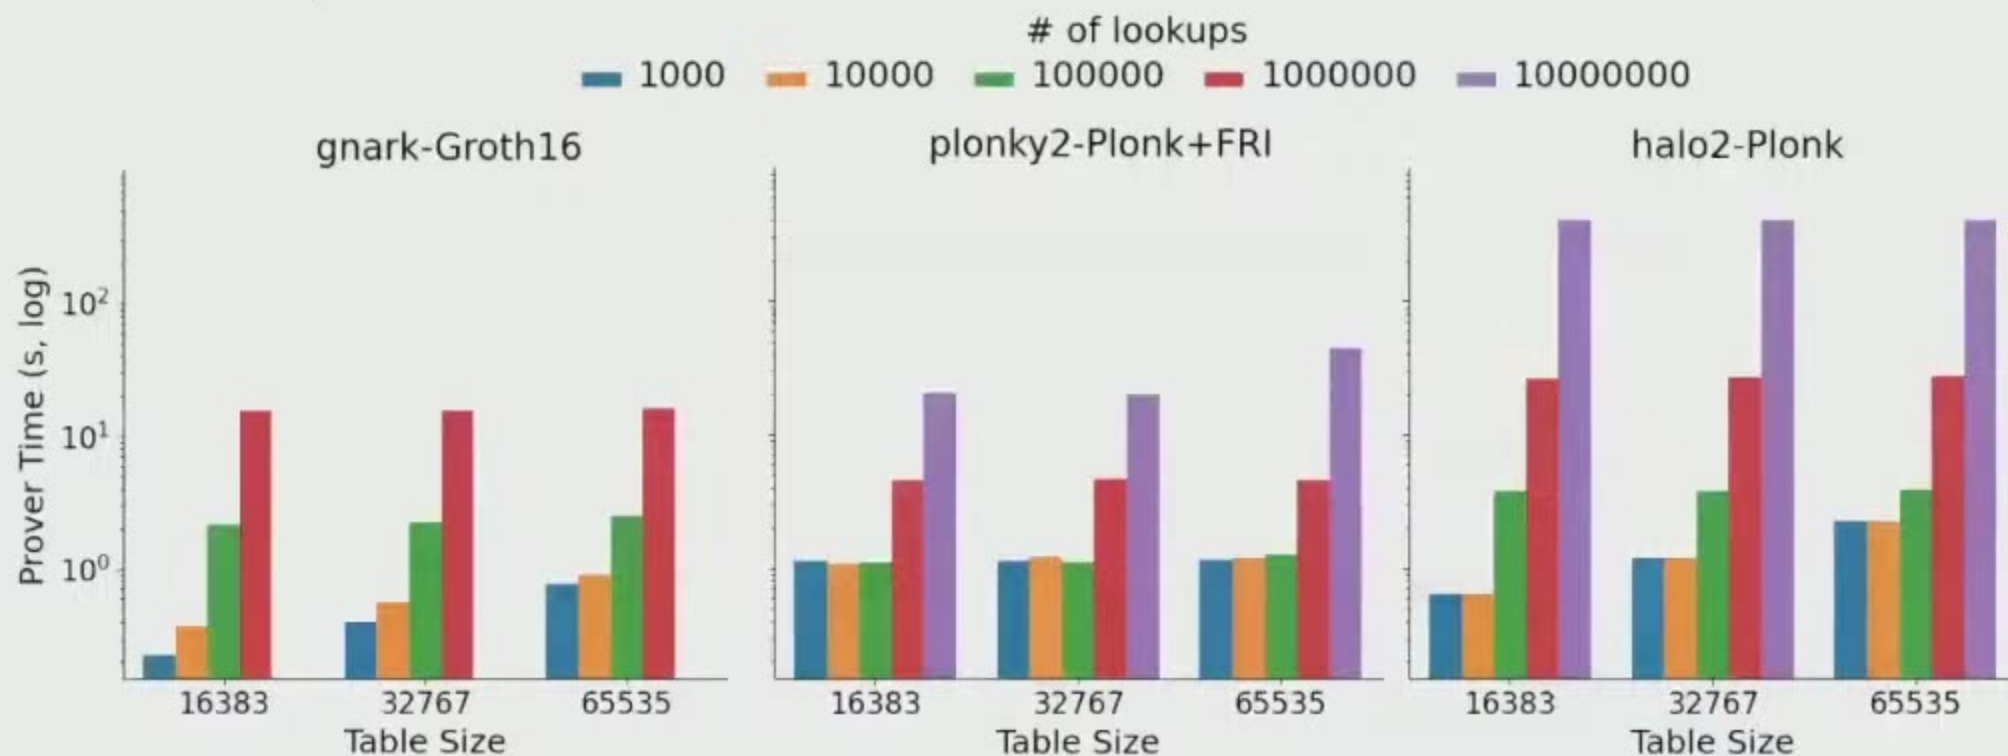
2. Output based on selector on sign bit

ReLU( $b$ ) lookup: commit table  
and add lookup constraints

<b>b</b>	<b>ReLU(b)</b>
-1	0
0	0
1	1
2	2



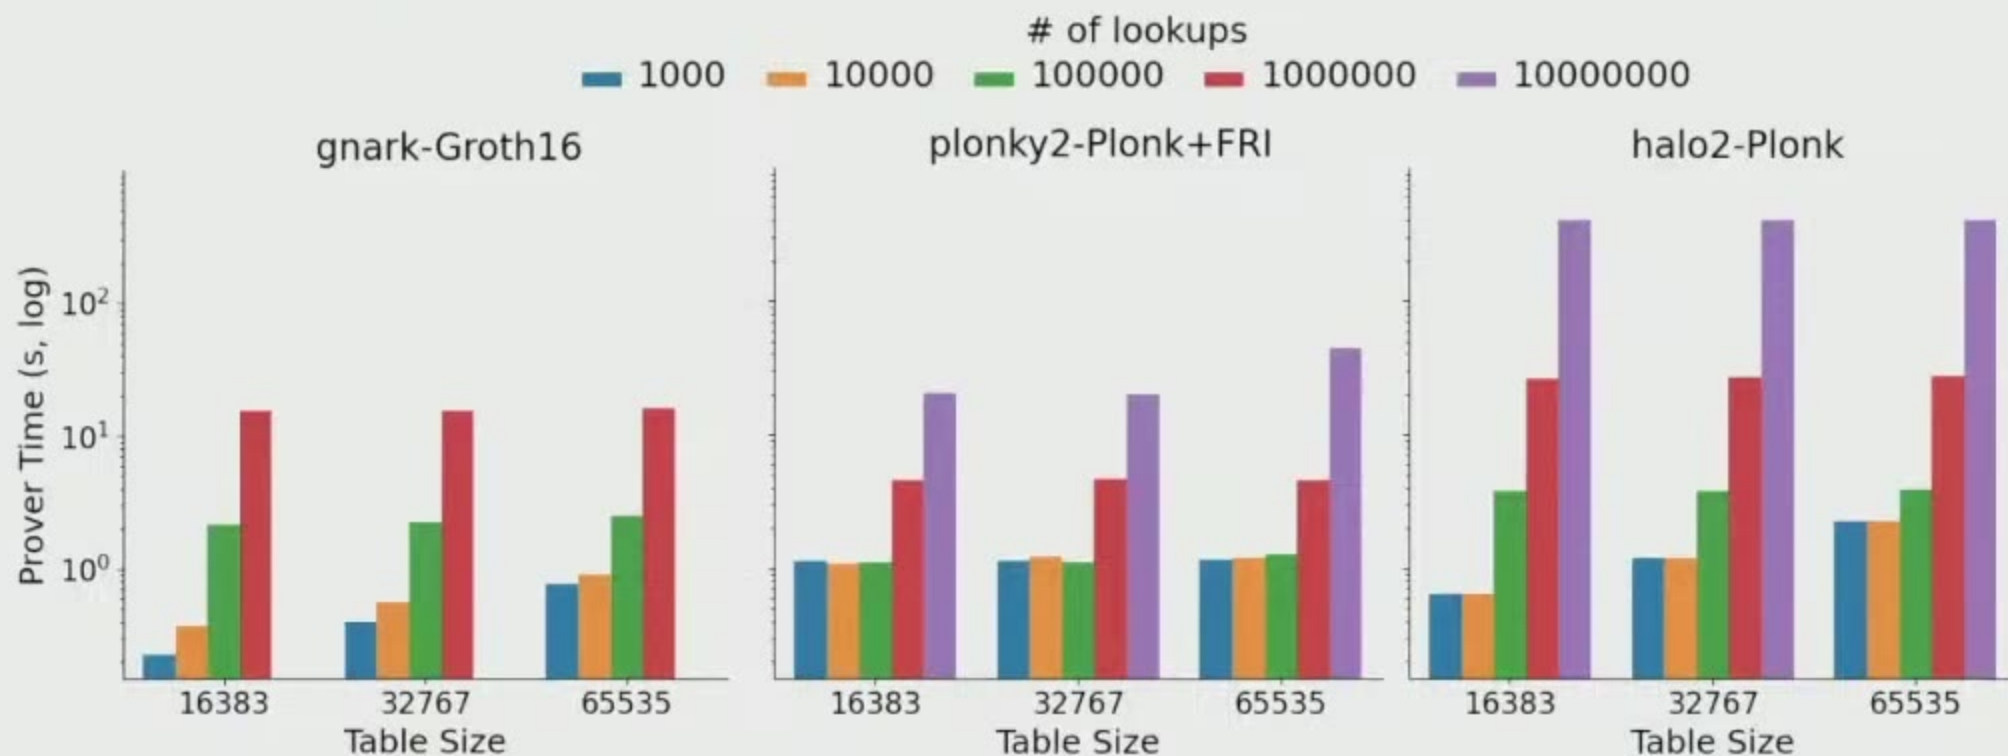
# Lookup costs



- Gnark's lookups are slow enough that bit-decomposition wins out
- Plonky2's lookups are most scalable



# Lookup costs

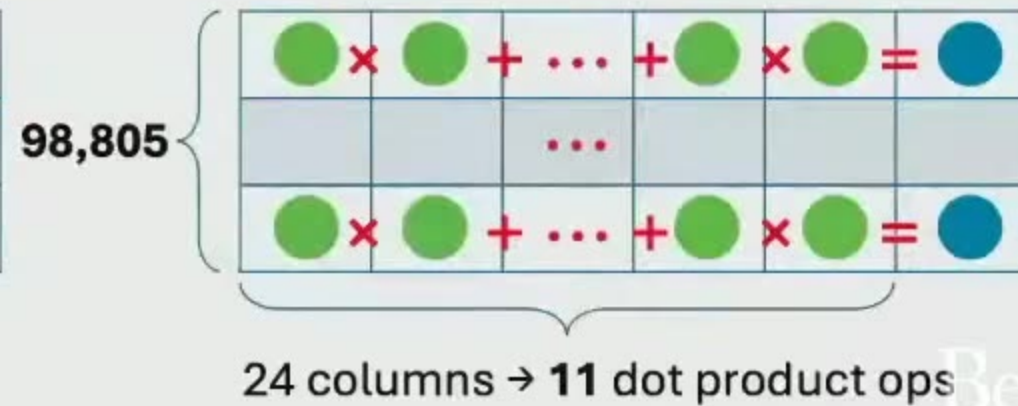
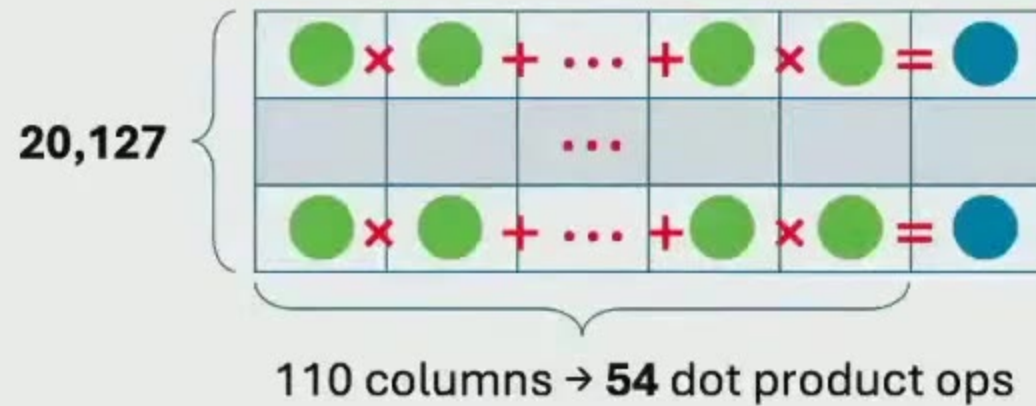


- Gnark's lookups are slow enough that bit-decomposition wins out
- Plonky2's lookups are most scalable

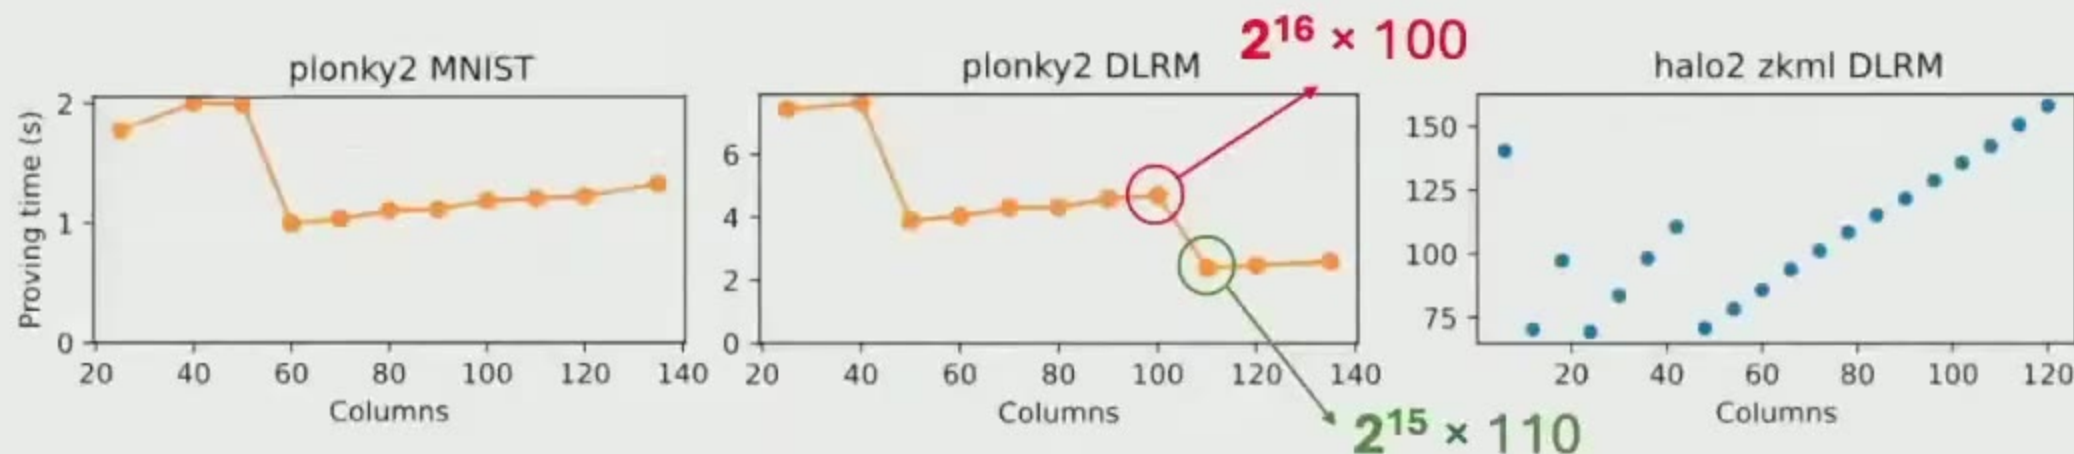
# Optimizations with circuit structure

- Halo2 and Plonky2 have rows and columns
- Rows affect sizes of operations: FFTs and MSMs
  - Padded to a power of 2
- Columns affect number of operations

Dot product gate configurations in DLRM



# Row and column microbenchmarks



Proving time drops when increasing columns decreases rows

# Conclusion

- ZKPerf: a ZK-SNARK proving benchmark with crypto and ML tasks
- Most costly parts of ML circuits are nonlinearities (e.g., ReLU)
  - Need efficient lookups to prove larger models
- Circuit structure matters
  - One way: optimize row and column dimensions



@liliatangxy



liliat2@illinois.edu



# Conclusion

- ZKPerf: a ZK-SNARK proving benchmark with crypto and ML tasks
- Most costly parts of ML circuits are nonlinearities (e.g., ReLU)
  - Need efficient lookups to prove larger models
- Circuit structure matters
  - One way: optimize row and column dimensions



@liliatangxy



liliat2@illinois.edu

# Conclusion

- ZKPerf: a ZK-SNARK proving benchmark with crypto and ML tasks
- Most costly parts of ML circuits are nonlinearities (e.g., ReLU)
  - Need efficient lookups to prove larger models
- Circuit structure matters
  - One-way optimizer and column dimensions



uiuc-kang-lab/zkperf



@liliatangxy



liliat2@illinois.edu



@liliatangxy



liliat2@illinois.edu



Berkeley  
UNIVERSITY OF CALIFORNIA

Powered by Zoom

# Conclusion

- ZKPerf: a ZK-SNARK proving benchmark with crypto and ML tasks
- Most costly parts of ML circuits are nonlinearities (e.g., ReLU)
  - Need efficient lookups to prove larger models
- Circuit structure matters
  - One way: optimize row and column dimensions



uiuc-kang-lab/zkperf



@liliatangxy



liliat2@illinois.edu

